

## What is NoSQL Database?

NoSQL stands for "Not Only SQL".

It is a type of database designed to **store, retrieve, and manage large volumes of unstructured, semi-structured, or rapidly changing data** that traditional **RDBMS (Relational Database Management Systems)** cannot efficiently handle.

Unlike relational databases, NoSQL databases **do not require fixed table schemas** and **can scale horizontally** across many servers.

## Why NoSQL is Required

Traditional **RDBMS** (like MySQL, Oracle, PostgreSQL) work well for structured data and fixed schema.

However, with the rise of **Big Data, IoT, social media, and cloud computing**, new challenges emerged:

Problem in RDBMS	Solution in NoSQL
Rigid schema (tables, columns)	Schema-less or dynamic schema
Difficult horizontal scaling	Easily scalable across servers
Performance issues with huge data	High-speed read/write operations
Not suitable for unstructured data (e.g. images, JSON, videos)	Supports unstructured/semi-structured data
Complex joins slow down performance	Data is often stored in a denormalized way for faster access

## Types of NoSQL Databases

### 1. Document-oriented Databases

- Store data as **JSON, BSON, or XML documents**.
- Each document contains all the information for a record.
- No need for predefined schema.

**Examples:** MongoDB, CouchDB, Firebase Firestore

**Use Case:** Content management systems, e-commerce product catalogs.

**Example Document:**

```
{  
    "student_id": 101,  
    "name": "Amit Kumar",  
    "courses": ["DBMS", "AI", "Python"],  
    "marks": {"DBMS": 85, "AI": 90, "Python": 92}  
}
```

## 2. Key-Value Stores

Simplest type of NoSQL.

Data stored as key-value pairs (like a dictionary or hash map).

Extremely fast for lookups.

Examples: Redis, DynamoDB, Riak

Use Case: Caching, session management, real-time recommendations.

Example:

```
"user101" : {"name:'Amit', age:22, city:'Pune'}"
```

## 3. Column-Family Stores

Data stored in columns instead of rows.

Designed for large analytical workloads and high read/write throughput.

Examples: Apache Cassandra, HBase, ScyllaDB

Use Case: Time-series data, IoT data storage, analytics systems.

Structure Example:

```
Row Key: Student101 Columns: {Name: Amit, Course: DBMS, Marks: 85}
```

## 4. Graph Databases

Store data as nodes (entities) and edges (relationships).

Ideal for representing complex relationships and networked data.

Examples: Neo4j, Amazon Neptune, ArangoDB

**Use Case:** Social networks, fraud detection, recommendation engines.

Example:

(Amit) - [FRIENDS\_WITH] -> (Rohit)

## Use Cases of NoSQL Databases

Domain	Example Use Case	NoSQL Type
Social Media	User profiles, connections	Graph DB
E-commerce	Product catalog, reviews	Document DB
Banking	Fraud detection, transaction graph	Graph DB
IoT Systems	Sensor data storage	Column-family DB
Gaming	Real-time leaderboard, session data	Key-value DB
Big Data Analytics	High-speed data ingestion	Column-family DB

## Key Benefits of NoSQL Databases

- Scalability:** Easily scaled horizontally across servers.
- Flexibility:** Schema-less design supports dynamic data.
- Performance:** Faster reads/writes for large datasets.
- High Availability:** Distributed architecture ensures uptime.
- Supports Modern Data:** Handles JSON, images, videos, logs, etc.

Difference between RDBMS and NoSQL

Feature	RDBMS	NoSQL
Data Model	Structured (tables)	Unstructured/Semi-structured
Schema	Fixed	Dynamic
Scalability	Vertical	Horizontal
Query Language	SQL	Non-SQL or API-based
Relationships	Supported (joins)	Limited or via graph
Examples	MySQL, Oracle	MongoDB, Cassandra, Redis, Neo4j

## Key Points on NoSQL

### 1. Do NoSQL Databases Have Primary Keys?

Yes, but not in the same way as RDBMS.

Every NoSQL database has a **way to uniquely identify each record** — though it might not be called a “primary key” formally.

NoSQL Type	Equivalent of Primary Key	Example
Key-Value Store	The <b>key</b> itself is the primary key.	"user101" → {name:'Amit'}
Document Store	Each document has a unique <b>_id</b> field.	{ "_id": 1, "name": "Amit" }
Column-Family Store	Each row has a <b>row key</b> (acts like primary key).	RowKey: student101
Graph Database	Each <b>node</b> has a unique identifier (node ID).	(user:101)

### 2. Do NoSQL Databases Have Foreign Keys?

No, most NoSQL databases do NOT use foreign keys.

- There are **no foreign key constraints** between collections or tables.
- Relationships are usually **embedded** or **referenced manually** inside documents.
- The database **does not enforce** referential integrity like an RDBMS.

Example in MongoDB (Document DB):

Instead of two tables with a foreign key relation, we might embed data:

```
{  
  "_id": 101,  
  "name": "Amit",  
  "courses": [  
    {"course_id": "DBMS", "marks": 85},  
    {"course_id": "AI", "marks": 90}  
  ]  
}
```

### 3. Do NoSQL Databases Support Joins?

Generally, NoSQL databases do not support joins like SQL does.

Why?

- Joins slow down performance in large distributed systems.
- NoSQL prefers **denormalized data**, meaning related information is often **stored together** in one document or collection.

However, there are **exceptions**:

- **MongoDB** supports a limited \$lookup operator that works like a left outer join.

- **Cassandra** and **Redis** generally avoid joins altogether.
- **Graph Databases (like Neo4j)** do not need joins — relationships are built-in and directly traversed.

```
db.orders.aggregate([
  {
    $lookup: {
      from: "customers",
      localField: "customer_id",
      foreignField: "_id",
      as: "customer_info"
    }
  }
]);
```

## Summary of Features

Feature	RDBMS	NoSQL
<b>Primary Key</b>	Yes	Yes (in all types, but implemented differently)
<b>Foreign Key</b>	Yes	Rarely used / Not enforced
<b>Joins</b>	Supported	Not supported or limited
<b>Data Relation Handling</b>	Through normalization and joins	Through embedding or linking

Q1. Show a small practical example in MongoDB comparing how the same data (Students & Courses) is stored with joins in SQL vs embedded in NoSQL ?

1. **RDBMS (MySQL)** – using primary key, foreign key, and joins
2. **NoSQL (MongoDB)** – using embedded documents (no joins)

## Example Scenario

We have two entities:

- **Students** – each student has an ID, name, and department.
- **Courses** – each student can take multiple courses with marks.

### 1. In RDBMS (MySQL)

#### ► Tables and Relationships

We use two tables with a foreign key.

```

-- Students Table
CREATE TABLE Students (
    student_id INT PRIMARY KEY,
    name VARCHAR(50),
    department VARCHAR(50)
);

-- Courses Table
CREATE TABLE Courses (
    course_id INT AUTO_INCREMENT PRIMARY KEY,
    student_id INT,
    course_name VARCHAR(50),
    marks INT,
    FOREIGN KEY (student_id) REFERENCES Students (student_id)
);

INSERT INTO Students VALUES ( 101, 'Amit Kumar', 'Computer Science');

INSERT INTO
    Students
VALUES (
    102,
    'Rohit Sharma',
    'Information Tech'
);

INSERT INTO
    Courses (
        student_id,
        course_name,
        marks
    )
VALUES (101, 'DBMS', 85),
       (101, 'AI', 90),
       (102, 'Networking', 88),
       (102, 'Python', 92);

```

## 2. In NoSQL (MongoDB)

In MongoDB, we typically embed the courses inside `each` student `document` – no `foreign key`, no `join`.

`Collection: students`

```
db.students.insertMany([
  {
    _id: 101,
    name: "Amit Kumar",
    department: "Computer Science",
    courses: [
      { course_name: "DBMS", marks: 85 },
      { course_name: "AI", marks: 90 }
    ]
  },
  {
    _id: 102,
    name: "Rohit Sharma",
    department: "Information Tech",
    courses: [
      { course_name: "Networking", marks: 88 },
      { course_name: "Python", marks: 92 }
    ]
  }
]);
```

► Query to fetch all students with their courses:

```
db.students.find({}, { name: 1, courses: 1, _id: 0 }).pretty();
```

Result (directly from single collection):

```
{
  "name": "Amit Kumar",
  "courses": [
    {"course_name": "DBMS", "marks": 85},
    {"course_name": "AI", "marks": 90}
  ]
}
{
  "name": "Rohit Sharma",
  "courses": [
    {"course_name": "Networking", "marks": 88},
    {"course_name": "Python", "marks": 92}
  ]
}
```

No JOIN needed – all related data is inside one document.

## Comparison Summary

Feature	RDBMS (MySQL)	NoSQL (MongoDB)
Data Model	Tables (rows & columns)	Documents (JSON-like)
Relationship	Separate tables linked by foreign key	Embedded sub-documents
Joins	Required for related data	Not needed
Schema	Fixed (predefined columns)	Flexible (dynamic)
Performance (large scale)	Slower due to joins	Faster (data retrieved in one read)

## How data is stored in JSON in NoSQL

In NoSQL databases (especially Document-oriented databases like MongoDB, CouchDB, and Firebase), data is stored as documents in a JSON-like format.

## What is JSON?

JSON (JavaScript Object Notation) is a text-based format used to represent structured data. It uses key-value pairs and nested objects/arrays, for example:

In NoSQL, each document like this is stored in a collection (similar to a table in RDBMS).

So in MongoDB:

- The database contains collections.
- Each collection contains documents.
- Each document is stored internally in a binary form of JSON (called BSON).

## Do we type JSON data manually?

Sometimes yes, but mostly no.

Here's what that means:

### ► Option 1: Manually (for small data or examples)

You can directly insert JSON manually in MongoDB (for testing, labs, or examples):

```
db.students.insertOne({  
    student_id: 101,  
    name: "Amit Kumar",  
    department: "Computer Science",  
    courses: [  
        { course_name: "DBMS", marks: 85 },  
        { course_name: "AI", marks: 90 }  
    ]  
}
```

```
];
});
```

MongoDB automatically stores this **JSON document** internally **as** **BSON (binary JSON)**.

► **Option 2:** Automatically (**in real** applications)

In **real**-world systems, **data is not** typed manually.

It **is** automatically **generated or sent to the database by an application** (like a website, Python **program**, or API).

Example **using** Python (PyMongo):

```
import pymongo

# connect to MongoDB
client = pymongo.MongoClient("mongodb://localhost:27017/")
db = client["college"]
students = db["students"]

# insert data from your application
data = {
    "student_id": 102,
    "name": "Rohit Sharma",
    "department": "Information Tech",
    "courses": [
        {"course_name": "Networking", "marks": 88},
        {"course_name": "Python", "marks": 92}
    ]
}
students.insert_one(data)
```

Here, the **program** sends a Python **dictionary**, which **is** automatically converted **to** **JSON** format when **stored in** MongoDB.

How **JSON data is stored** internally

Steps:

- 1 You **insert JSON (or similar data)** called **JSON text**
- 2 Database stores it internally **BSON (Binary JSON)** faster, supports extra **data types**
- 3 When you retrieve it Returned back as **JSON**

Why JSON format is used in NoSQL

Advantage	Explanation
<b>Flexible Schema</b>	You can add or remove fields easily without altering the structure.
<b>Nested Data</b>	You can store arrays and objects (like courses inside students).
<b>Human-Readable</b>	Easy to understand and debug.
<b>Compatible</b>	Works well with APIs, web apps, and programming languages.

Give me a Real-World Example: [Data Flow in NoSQL](#)

Example: Ride-Sharing App (like Uber or Ola)

Every second, thousands of events happen:

- New rides requested
- Driver location updates (GPS)
- Payment logs
- Ratings and feedback

All this data is:

- Generated rapidly by mobile devices.
- Sent to a backend server (API).
- Stored in MongoDB almost instantly as JSON-like documents.

#### Advantage

- Flexible Schema You can add or remove fields easily without altering the structure.
- Nested Data You can store arrays and objects (like courses inside students).
- Human-Readable, Easy to understand and debug.
- Compatible, Works well with APIs, web apps, and programming languages.

```
{  
  "ride_id": "R456",  
  "user_id": "U101",  
  "driver_id": "D55",
```

```

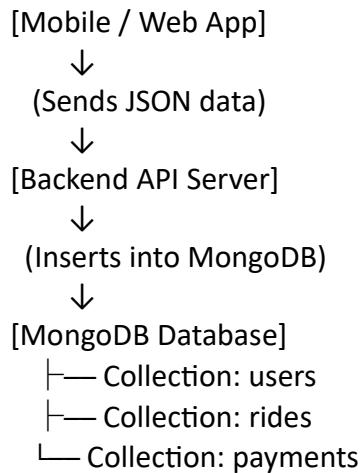
    "pickup": "Pune Station",
    "drop": "Airport",
    "start_time": "2025-11-07T10:15:00",
    "status": "completed",
    "fare": 320.50,
    "rating": 5
}

```

## Where MongoDB Excels (Use Cases)

Application Type	Why MongoDB Fits
E-commerce websites	Product details vary (different fields, reviews, images). JSON allows flexible storage.
Social media apps	Posts, comments, likes — all dynamic and fast-changing data.
IoT systems	Billions of sensor readings per second — stored rapidly.
Gaming apps	Real-time player stats, leaderboards, and game events.
Analytics dashboards	Continuous log and event data ingestion.

## Example Architecture



### MongoDB is ideal when:

- Data comes rapidly or in huge volumes.
- The data format changes often (dynamic fields).
- You need fast writes and reads.
- You want to scale easily across multiple servers.